

■■■ Q-MoleGen

Quantum-Enhanced Molecular Generation & Multi-Objective Property Discovery Platform

A Hybrid Classical-Quantum Machine Learning Framework for *In Silico* De Novo Candidate Discovery, Delaney ESOL Solubility Prediction, and Pareto Frontier Optimization

PROJECT GUIDE & SUPERVISOR

Dr. G. JayaSuma

Professor, Department of Information Technology

Email: jayasuma@jntuk.edu.in

STUDENT RESEARCH TEAM & CONTRIBUTORS

- **D. Pravallika** (Reg. No: 22031A1215)
- **S. Sai Pavan** (Reg. No: 22031A1255)
- **G. Vyshnavi** (Reg. No: 22031A1219)
- **T. Jagadeesh** (Reg. No: 22031A1260)

OPEN-SOURCE REPOSITORY & LIVE PRODUCTION LINKS

- **GitHub Repository:** <https://github.com/VALLI-0114/Q-MoleGen>
- **Live Vercel Deployment:** <https://q-mole-gen-git-main-kundum-pravallikas-projects.vercel.app>
- **Database Infrastructure:** Supabase PostgreSQL 15 Cloud Database Cluster

EXECUTIVE ABSTRACT: Traditional *de novo* molecular design suffers from massive candidate attrition, high screening costs (\$2.6B+ per drug), and intractable chemical search spaces ($\sim 10^{60}$ structures). **Q-MoleGen** introduces a unified, full-stack hybrid computing system that accelerates lead generation by synergizing: (1) RDKit cheminformatics filters, (2) classical ensemble regressors (Gradient Boosting, Random Forest) achieving **94.25% test accuracy** on Delaney ESOL solubility benchmarks, (3) Parameterized Quantum Support Vector Classifiers (QSVC) executing 4-qubit ZZFeatureMap circuits in a 16-dimensional Hilbert space, (4) multi-objective Pareto frontier ranking (solubility LogS, QED drug-likeness, SAS synthetic accessibility), and (5) a high-aesthetic, mobile-responsive React 18 / Vite web application connected to Supabase PostgreSQL cloud storage.

1. Problem Statement & Scientific Motivation

Small-molecule drug discovery requires searching an astronomically large chemical space estimated at **10^{60} feasible drug-like compounds**. Conventional wet-lab High-Throughput Screening (HTS) evaluates molecules empirically, requiring 10–15 years and over \$2.6 billion per commercially approved drug with an overall clinical attrition rate exceeding 90%.

A critical reason for candidate failure during early preclinical screening is poor **physico-chemical properties**, predominantly insufficient **aqueous solubility (LogS)**, unfavorable **lipophilicity (LogP)**, high synthetic complexity, and poor oral bioavailability (Lipinski Rule of Five violations).

Key Bottlenecks Addressed by Q-MoleGen:

- **Combinatorial Search Inefficiency:** Classical brute-force screening cannot scale to high-dimensional chemical spaces without intelligent generative filtering.
- **Single-Objective Sub-optimality:** Optimizing only for potency frequently creates highly insoluble or synthetically inaccessible molecules. Multi-objective Pareto optimization is mandatory.
- **Quantum Mechanical Complexity:** Molecular bonding and electron correlation are fundamentally quantum phenomena. Quantum Machine Learning (QML) kernels enable non-linear Hilbert space mappings inaccessible to standard linear kernels.

2. Comprehensive Technology & Software Stack

Each core technology layer was selected to optimize performance, computational fidelity, and user experience across the discovery pipeline:

Layer / Module	Technology & Version	Core Role & Capabilities in Q-MoleGen
Cheminformatics	RDKit (2024.x)	Molecular graph parsing, SMILES canonicalization, 2D/3D descriptor extraction, Morgan ECFP4 fingerprinting, QED and SAS computation.
Quantum Computing	Qiskit 2.x & Aer	Parameterized Quantum Circuits (PQC), 4-qubit ZZFeatureMap ansatz, 16-dim Hilbert state evaluation, and quantum fidelity kernel computation.
Classical ML	Scikit-Learn (1.4+)	Gradient Boosting, Random Forest, RBF SVM, Linear SVM, and Ridge regression pipelines; cross-validation and hyperparameter tuning.
Deep Learning	PyTorch (2.2+)	Generative candidate generation tensors, neural representation encoders, and gradient-based property surrogates.
Backend Server	Django 5.0 + REST	Secure REST APIs, JSON serialization, user session registry, experiment coordination, and error logging.

Frontend UI/UX	React 18 & Vite 5	Single-Page Application (SPA) with Outfit & Plus Jakarta Sans typography, responsive mobile drawer, and glassmorphism design system.
Data Visualization	Chart.js & React-Chartjs-2	Dynamic interactive property histograms, Lipinski Ro5 doughnut charts, Pareto scatter plots, and generation progress curves.
Cloud Persistence	Supabase (PostgreSQL 15)	Cloud database synchronization for user accounts, generated molecule libraries, experiment campaigns, and inquiries.
Production Hosting	Vercel Platform	Continuous integration and global edge CDN deployment with SPA client-side routing.

3. Dataset Architecture & Benchmark Corpora

Q-MoleGen utilizes standard, peer-reviewed physical chemistry and quantum benchmark datasets:

3.1 Delaney ESOL Aqueous Solubility Dataset

The primary benchmark used is the **Delaney ESOL dataset** (1,128 compounds), capturing experimentally verified aqueous solubility measurements ($\log S$ in $\text{ext}\{\text{mols/L}\}$). The dataset spans a wide dynamic solubility range from -11.6 (extremely insoluble) to $+1.58$ (highly soluble), representing diverse pharmaceutical scaffolds including aromatics, heterocycles, aliphatics, and polycyclics.

Data Preprocessing Pipeline:

- SMILES Canonicalization:** Standardized stereochemistry, neutralized salts, and stripped counter-ions.
- RDKit Sanity Check:** Verified valence validity and bond topologies (100% of Delaney passed).
- Stratified Train/Test Split:** 80% training set (902 compounds), 20% independent test set (226 compounds).

3.2 QM9 Quantum Chemistry Dataset

The **QM9 dataset** contains 133,885 small organic molecules (up to 9 heavy atoms: C, N, O, F) with geometric, energetic, and electronic properties computed via Density Functional Theory (DFT) at the B3LYP/6-31G(2df,p) level. Q-MoleGen extracts HOMO-LUMO energy gaps ($\Delta\epsilon$), dipole moments (μ), and polarizabilities to evaluate quantum electronic states.

4. Molecular Feature Engineering & Representations

To enable both classical and quantum algorithms to process chemical structures, Q-MoleGen transforms 1D SMILES strings into comprehensive descriptor feature vectors:

Feature Type	Dimension	Extracted Properties & Algorithms	Impact on Solubility (LogS)
2D Physico-chemical	6 Descriptors	Molecular Weight (MW), LogP (Wildman-Crippen), TPSA, H-Bond Donors (HBD), H-Bond Acceptors (HBA), Rotatable Bonds.	Strong inverse correlation between MW/LogP and LogS; positive correlation with TPSA.
Morgan Fingerprints (ECFP4)	1024 / 2048 Bits	Circular topological fingerprints with radius 2 capturing circular atom environments and functional groups.	Captures specific hydrophilic/hydrophobic fragment contributions (hydroxyl, carboxyl, aromatic).
Quantum State Mapping	4 Qubits (16-Dim)	Continuous features mapped to qubit rotation angles via parameterized ZZFeatureMap.	Enables quantum Hilbert space distance kernel evaluation.

5. Classical Machine Learning & Predictive Modeling

Six distinct supervised learning algorithms were benchmarked on Delaney ESOL test data to establish baseline predictive performance for solubility estimation:

Model Architecture	Test Accuracy	Precision	Recall	F1-Score	ROC-AUC	Fit Time
Gradient Boosting (Champion)	94.25%	94.8%	94.0%	94.4%	0.977	0.133s
Random Forest Regressor	93.36%	92.5%	94.9%	93.7%	0.975	0.117s
QSVC (ZZ-FeatureMap 4-Qubit)	89.82%	89.8%	90.6%	90.2%	0.959	0.793s
Support Vector Regressor (RBF)	89.38%	88.4%	91.5%	89.9%	0.964	0.033s
Linear Support Vector Regressor	88.50%	86.4%	92.3%	89.3%	0.946	0.021s
Logistic Regression Baseline	88.05%	85.7%	92.3%	88.9%	0.944	0.005s

Key Finding: Gradient Boosting emerged as the champion classical ensemble, achieving **94.25% test accuracy** and a **0.977 ROC-AUC** with high compute efficiency (0.133s training time). Random Forest followed closely with 93.36% accuracy.

6. Quantum Machine Learning (QML) & Circuit Topology

To evaluate potential quantum advantage in non-linear chemical classification, Q-MoleGen implements a **Quantum Support Vector Classifier (QSVC)** utilizing Qiskit's parameterized quantum kernel estimation.

6.1 4-Qubit ZZ-FeatureMap Circuit Architecture

The feature map maps a classical feature vector $\mathbf{x} \in \mathbb{R}^4$ (normalized MW, LogP, TPSA, and HBD) into a 16-dimensional quantum state $|\Phi(\mathbf{x})\rangle \in \mathbb{C}^{16}$ across 4 qubits with a total circuit depth of 19 gates:

- **Layer 1 — Superposition:** Hadamard gates $H^{\otimes 4}$ applied to all 4 qubits to create an equal superposition state $|+\rangle^{\otimes 4}$.
- **Layer 2 — Single-Qubit Phase Encoding:** Single-qubit phase rotation gates $P(2x_i)$ encoding individual descriptor magnitudes.
- **Layer 3 — Cross-Qubit Entanglement:** Controlled-Phase interactions $CX_{\{(i,j)\}} \cdot P(2(\pi - x_i)(\pi - x_j)) \cdot CX_{\{(i,j)\}}$ creating multi-body quantum entanglement between all pairs of qubits (q_0, q_1, q_2, q_3) .

6.2 Quantum Kernel Inner Product Evaluation

The quantum kernel entry between molecules $\mathbf{x}_i$ and $\mathbf{x}_j$ is computed as the transition probability (fidelity overlap) of their prepared quantum states:

$$K(\mathbf{x}_i, \mathbf{x}_j) = |\langle \Phi(\mathbf{x}_i) | \Phi(\mathbf{x}_j) \rangle|^2$$

The computed $N \times N$ quantum kernel matrix is then supplied to the dual SVM quadratic optimizer. On the Delaney ESOL benchmark, QSVC achieved **89.82% accuracy** and **0.959 ROC-AUC**, demonstrating robust quantum state representation.

7. De Novo Molecular Generation & Quality Validation

Q-MoleGen's generative engine generates candidate structures targeted to user-defined physical property ranges (e.g. High Solubility $\log S > -2.0$, High QED Drug-likeness > 0.60 , Low Synthetic Complexity $\text{ext}\{\text{SAS}\} \leq 3.5$).

7.1 Generative Quality Metrics (Empirical Batch #1)

Metric Name	Observed Value	Benchmark Target	Significance in Discovery
Total Generated	20 Molecules	20 Target	Candidate batch size for exploration campaign.
Validity Rate	100.0% (20/20)	> 95.0%	All candidates pass RDKit chemical valence and aromaticity sanity checks.
Uniqueness Rate	100.0% (20/20)	> 90.0%	Zero duplicate SMILES strings in the generated batch.
Novelty Rate	70.0% (14/20)	> 60.0%	14 candidate structures are entirely novel relative to the 1,128 Delaney reference library.
Lipinski Ro5 Compliance	100.0% (0 Violations)	> 80.0%	All 20 generated candidates satisfy $\text{MW} \leq 500$, $\text{LogP} \leq 5$, $\text{HBD} \leq 5$, $\text{HBA} \leq 10$.

8. Multi-Objective Pareto Optimization Engine

Candidate molecules in drug discovery must satisfy conflicting objectives: high potency, high solubility, acceptable permeability, and easy chemical synthesis. Q-MoleGen formulates a **Pareto Dominance Solver** to prioritize candidates on the non-dominated Pareto frontier.

8.1 Mathematical Pareto Formulation

A molecule $\mathbf{A}$ dominates molecule $\mathbf{B}$ ($\mathbf{A} \succ \mathbf{B}$) if and only if:

- $\forall i \in \{1, \dots, k\}, f_i(\mathbf{A}) \geq f_i(\mathbf{B})$ (A is at least as good as B in all objectives)
- $\exists j \in \{1, \dots, k\}, f_j(\mathbf{A}) > f_j(\mathbf{B})$ (A is strictly better than B in at least one objective)

Where objectives include: $f_1 = \text{ext}\{\text{LogS}\}$ (Solubility), $f_2 = \text{ext}\{\text{QED}\}$ (Drug-likeness), $f_3 = -\text{ext}\{\text{SAS}\}$ (Synthetic Accessibility).

8.2 Top Pareto-Optimal Candidates Identified

Candidate ID	Canonical SMILES	Predicted LogS	QED Score	SAS Score	Multi-Obj Score	Pareto Status
QMOL-001	<chem>O=C(O)c1cccc1O</chem>	-1.84 mol/L	0.77	1.34	78.1	Rank 1 (Non-dominated)
QMOL-002	<chem>CC(=O)Oc1cccc1C(=O)O</chem>	-2.26 mol/L	0.75	1.65	74.2	Rank 1 (Non-dominated)
QMOL-003	<chem>CC(C)c1ccc(C(C)C(=O)O)cc1</chem>	-2.91 mol/L	0.82	2.10	71.6	Rank 1 (Non-dominated)
QMOL-004	<chem>CN1C=NC2=C1C(=O)N(C(=O)N2C)C</chem>	-1.15 mol/L	0.68	1.98	69.8	Rank 1 (Non-dominated)
QMOL-005	<chem>CC(=O)Nc1ccc(O)cc1</chem>	-1.42 mol/L	0.71	1.42	68.4	Rank 1 (Non-dominated)

9. Full-Stack Integration & System Architecture

Q-MoleGen connects computational chemistry algorithms with modern cloud-native web engineering to create an accessible discovery workstation:

SYSTEM INTEGRATION WORKFLOW:

1. Client Presentation Layer (React 18 + Vite):

- **Researcher Portal:** Central command center showing live dynamic KPI counters (20 candidate leads, 16-dim Hilbert state, 94.25% baseline accuracy), 3D molecular canvas, and saved campaign tables with 1-click delete & CSV export.
- **Experiment Analytics:** Interactive dropdown selector with dynamic histograms, Lipinski Ro5 doughnut charts, generational progress curves, and classical vs quantum performance benchmark tables.
- **SMILES Inspector & Generator:** Instant chemical parsing, 2D/3D structure rendering, and property score validation.
- **Mobile Navigation Drawer:** Responsive slide-down navigation menu with touch controls on screens $\leq 960\text{px}$.

2. Backend Orchestration Layer (Django REST Framework):

- `/api/generate/` — Executes generative sampling and RDKit filtering pipelines.
- `/api/researcher/stats/` — Computes live summary KPIs directly from active datasets.
- `/api/analytics/data/<id>/` — Serves dynamic histogram distributions and benchmark comparisons.
- `/api/auth/register/` — Role-based persona gate strictly scoped to Researcher and Admin.

3. Cloud Persistence Layer (Supabase PostgreSQL 15):

- Tables for registered users, saved research campaigns, generated candidate libraries, and scientific inquiries.

10. Conclusion & Future Horizons

10.1 Key Achievements:

- Successfully integrated **RDKit cheminformatics**, **Qiskit 2.x quantum simulation**, **classical ML ensembles**, and **multi-objective Pareto optimization** into an end-to-end platform.
- Champion Gradient Boosting model attained **94.25% test accuracy** on Delaney ESOL, while the 4-qubit QSVC kernel demonstrated **89.82% quantum classification accuracy** in a 16-dimensional Hilbert space.
- Generated 20 validated candidates with **100% chemical validity**, **70% novelty**, and **100% Lipinski Rule of Five compliance**.
- Deployed a modern, responsive full-stack platform live on Vercel with Supabase cloud database integration.

10.2 Future Horizons:

- **Fault-Tolerant Quantum Hardware Execution:** Transitioning from Qiskit Aer simulators to physical IBM Quantum superconducting quantum processors (QPUs).
- **Active Learning & Automated Wet-Lab Loop:** Integrating robotic synthesis protocols with active learning Bayesian optimization.
- **3D Molecular Docking Integration:** Direct binding affinity scoring with target viral/oncology protein pockets using AutoDock Vina.

■ ■ Computational Drug-Likeness Heuristic Disclaimer:

The property predictions, LogS solubility values, Lipinski compliance metrics, and Pareto rankings generated by Q-MoleGen are in silico computational heuristics. They serve to prioritize candidates for synthesis and do NOT constitute clinical validation or proof of safety and pharmacological efficacy.